



# CKA 프로젝트

BSFAN | 김태경 박종승 윤재영 김효은

# 2개의 작업 클러스터 생성, 백업, 업그레이드

# 1. ETCD 백업

- 1. https://127.0.0.1:2379에서 실행 중인 etcd의 snapshot을 생성하고 snapshot을 /data/etcd-snapshot.db에 저장합니다.
- 2. 다음 다시 스냅샷을 복원합니다.
- 3. etcdctl을 사용하여 서버에 연결하기 위해 다음 TLS 인증서/키가 제공됩니다.
  - CA certificate: /etc/kubernetes/pki/etcd/ca.crt
  - Client certificate: /etc/kubernetes/pki/etcd/server.crt
  - Client key: /etc/kubernetes/pki/etcd/server.key

## etcd 백업

- 1. CA certificate: /etc/kubernetes/pki/etcd/ca.crt
- 2. Client certificate: /etc/kubernetes/pki/etcd/server.crt
- 3. Client key: /etc/kubernetes/pki/etcd/server.key

```
root@k8s-master:~# ETCDCTL_API=3 etcdctl --endpoints=https://127.0.0.1:2379 --cacert=/etc/kubernetes/pki/etcd/ca.crt etes/pki/etcd/server.crt --key=/etc/kubernetes/pki/etcd/server.key snapshot save /data/etcd-snapshot.db

[3]
```

etcd 복원

/data/etcd- snapshot.db

## 확인

# 2개의 작업 클러스터 생성, 백업, 업그레이드

# 2. Cluster Upgrade

- 1. 마스터 노드의 모든 Kubernetes control plane및 node 구성 요소를 버전 1.29.6-1.1 버전으로 업그레이드합니다.
- 2. master 노드를 업그레이드하기 전에 drain 하고 업그레이드 후에 uncordon해야 합니다. ["주의사항" 반드시 Master Node에서 root권한을 가지고 작업을 실행해야 한다.]

#### master drain

kubectl drain k8s-master --ignore-daemonsets

```
root@k8s-master:~# kubectl drain k8s-master --ignore-daemonsets
node/k8s-master already cordoned
Warning: ignoring DaemonSet-managed Pods: kube-system/kube-proxy-vcgrw, kube-system/weave-net-lchss
evicting pod kube-system/coredns-5dd5756b68-z4gmf
evicting pod kube-system/coredns-5dd5756b68-6fgkp
pod/coredns-5dd5756b68-z4gmf evicted
pod/coredns-5dd5756b68-6fgkp evicted
node/k8s-master drained
```

master uncordon

kubectl uncordon k8s-master

root@k8s-master:~# kubectl uncordon k8s-master node/k8s-master uncordoned

# 역할 기반 제어

## 3. Service Account, Role, RoleBinding 생성

애플리케이션 운영중 특정 namespace의 Pod들을 모니터할수 있는 서비스가 요청되었습니다. api-access 네임스페이스의 모든 pod를 view할 수 있도록 다음의 작업을 진행하시오.

- api-access라는 새로운 namespace에 pod-viewer라는 이름의 Service Account를 만듭니다.
- podreader-role이라는 이름의 Role과 podreader-rolebinding이라는 이름의 RoleBinding을 만듭니다.
- 앞서 생성한 ServiceAccount를 API resource Pod에 대하여 watch, list, get을 허용하도록 매핑하시오.

네임스페이스 생성

```
guru@k8s-master:~$ kubectl create ns api-access
namespace/api-access created
```

네임스페이스 확인

```
guru@k8s-master:~$ kubectl get ns api-access
NAME STATUS AGE
api-access Active 20m
```

Service Accont 생성

sa pod-viewer -n api-access

```
guru@k8s-master:~$ kubectl create sa pod-viewer -n api-access
serviceaccount/pod-viewer created
```

Service Accont 확인

kubectl get sa -n api-access

```
guru@k8s-master:~$ kubectl get sa -n api-access
NAME SECRETS AGE
default 0 4m59s
pod-viewer 0 18s
```

Role 생성 - name: podreader-role / resource pod에 대하여 watch, list, get 허용

kubectl create role podreader-role --resource=pod --verb=watch,list,get -n api-access

```
guru@k8s-master:~$ kubectl create role podreader-role --resource=pod --verb=watch,list,get, -n api-access
Warning: '' is not a standard resource verb
role.rbac.authorization.k8s.io/podreader-role created
```

kubectl describe role -n api-access

RoleBinding 생성 - name: podreader-rolebinding

kubectl create rolebinding podreader-rolebinding --role=podreader-role --serviceaccount=api-access:pod-viewer -n api-access

```
guru@k8s-master:~$ kubectl create rolebinding podreader-rolebinding --role=podreader-role --serviceaccount=api-acces s:pod-viewer -n api-access \rolebinding.rbac.authorization.k8s.io/podreader-rolebinding created
```

#### Rolebindind 확인

kubectl describe rolebinding -n api-access

# 역할 기반 제어

## 4. Service Account, ClusterRole, ClusterRoleBinding 생성

애플리케이션 배포를 위해 새로운 ClusterRole을 생성하고 특정 namespace의 ServiceAccount를 바인드하시오.

- 1. 다음의 resource type에서만 Create가 허용된 ClusterRole deployment-clusterrole을 생성합니다.
  - Resource Type: Deployment StatefulSet DaemonSet
- 2. 미리 생성된 namespace api-access 에 cicd-token이라는 새로운 ServiceAccount를 만듭니다.
- 3. ClusterRole deployment-clusterrole을 namespace api-access 로 제한된 새 ServiceAccount cicd-token 에 바인딩하세요.
- 2. 네임스페이스 api-access에 cicd-token 새로운 sa 생성

```
guru@k8s-master:~$ kubectl create sa cicd-token -n api-access serviceaccount/cicd-token created
```

1. resource type에서만 Create가 허용된 ClusterRole deployment-clusterrole을 생성

```
guru@k8s-master:~$ kubectl create clusterrole deployment-clusterrole --resource=deployment,statefulset,daemonset --v
erb=create
clusterrole.rbac.authorization.k8s.io/deployment-clusterrole created
```

3. ClusterRole deployment-clusterrole을 namespace api-access 로 제한된 새 ServiceAccount cicd-token에 바인딩

```
guru@k8s-master:~$ kubectl create clusterrolebinding deployment-clusterrolebinding --serviceaccount=api-access:cicd-token --clusterrole=deployment-clusterrole -n api-access clusterrolebinding.rbac.authorization.k8s.io/deployment-clusterrolebinding created
```

## 확인

Service Account

kubectl describe sa -n api-access cicd-token

```
guru@k8s-master:~$ kubectl describe sa -n api-access cicd-token
Name:
                      cicd-token
Namespace:
                      api-access
Labels:
                      <none>
Annotations:
                      <none>
Image pull secrets:
                      <none>
Mountable secrets:
                      <none>
Tokens:
                      <none>
Events:
                      <none>
guru@k8s-master:~$
```

#### ClusterRole

kubectl describe -n api-access clusterrole deployment-clusterrole

```
guru@k8s-master:~$ kubectl describe -n api-access clusterrole deployment-clusterrole
Name:
              deployment-clusterrole
Labels:
              <none>
Annotations:
              <none>
PolicyRule:
                                                          Verbs
  Resources
                      Non-Resource URLs
                                         Resource Names
                                         daemonsets.apps
                      [\ ]
                                                          [create]
                      deployments.apps
                                                          [create]
  statefulsets.apps
                                                          [create]
```

#### 확인

#### ClusterRole Binding

kubectl describe -n api-access clusterrolebinding deployment-clusterrolebinding

```
guru@k8s-master:~$ kubectl describe -n api-access clusterrolebinding deployment-clusterrolebinding
Name:
              deployment-clusterrolebinding
Labels:
              <none>
Annotations: <none>
Role:
        ClusterRole
  Kind:
  Name:
         deployment-clusterrole
Subjects:
  Kind
                  Name
                              Namespace
  ServiceAccount cicd-token api-access
```

# 노드관리

## 5. 노드 비우기

k8s-worker2 노드를 스케줄링 불가능하게 설정하고, 해당 노드에서 실행 중인 모든 Pod을 다른 node로 reschedule 하세요.

## node 확인

#### kubectl get no

```
guru@k8s-master:~$ kubectl get no
NAME
               STATUS
                         ROLES
                                          AGE
                                                 VERSION
k8s-master
               Ready
                         control-plane
                                          293d
                                                 v1.28.8
k8s-worker1
                                          293d
                                                 v1.28.8
               Ready
                         <none>
k8s-worker2
               Ready
                                          293d
                                                 v1.28.8
                         <none>
```

테스트 deploy 파일 생성

kubectl create deploy deploy-test --image=nginx:1.14 --replicas=2

guru@k8s-master:~\$ kubectl create deploy deploy-test --image=nginx:1.14 --replicas=2 deployment.apps/deploy-test created

kubectl get po -o wide

```
guru@k8s-master:~$ kubectl get
NAME
                                 READY
                                          STATUS
                                                    RESTARTS
                                                                AGE
                                                                       ΙP
                                                                                    NODE
                                                                                                   NOMINATED NODE
                                                                                                                     READINE
SS GATES
deploy-test-5cd79d8497-fmj72
                                 1/1
                                                                       10.32.0.3
                                                                                    k8s-worker2
                                          Running
                                                    Θ
                                                                46s
                                                                                                   <none>
                                                                                                                     <none>
deploy-test-5cd79d8497-gc7r2
                                 1/1
                                          Running
                                                    0
                                                                46s
                                                                       10.32.0.2
                                                                                   k8s-worker2
                                                                                                   <none>
                                                                                                                     <none>
```

k8s-worker2 노드 스케줄링 Disabled 설정 및 확인

kubectl cordon k8s-worker2

kubectl get no

```
guru@k8s-master:~$ kubectl cordon k8s-worker2
node/k8s-worker2 cordoned
guru@k8s-master:~$ kubectl get no
NAME
               STATUS
                                                             AGE
                                           ROLES
                                                                    VERSION
k8s-master
               Ready
                                            control-plane
                                                             293d
                                                                    v1.28.8
k8s-workerl
               Ready
                                                             293d
                                                                    v1.28.8
                                            <none>
               Ready, SchedulingDisabled
k8s-worker2
                                                             293d
                                                                    v1.28.8
                                           <none>
```

node drain

kubectl drain k8s-worker2 --ignore-daemonsets

```
guru@k8s-master:~$ kubectl drain k8s-worker2 --ignore-daemonsets
node/k8s-worker2 already cordoned
Warning: ignoring DaemonSet-managed Pods: kube-system/kube-proxy-7rl68, kube-system/weave-net-sbnc5
evicting pod default/deploy-test-5cd79d8497-gc7r2
evicting pod default/deploy-test-5cd79d8497-fmj72
pod/deploy-test-5cd79d8497-gc7r2 evicted
pod/deploy-test-5cd79d8497-fmj72 evicted
node/k8s-worker2 drained
```

확인

kubectl get po -o wide1

kubectl get no

```
guru@k8s-master:~$ kubectl get
                                 po -o wide
NAME
                                 READY
                                         STATUS
                                                    RESTARTS
                                                                AGE
                                                                        ΤP
                                                                                     NODE
                                                                                                    NOMINATED NODE
                                                                                                                      READT
NESS GATES
deploy-test-5cd79d8497-7vtt5
                                                                        10.40.0.2
                                                                                    k8s-worker1
                                 1/1
                                         Runnina
                                                    0
                                                                2m10s
                                                                                                    <none>
                                                                                                                      <none
deploy-test-5cd79d8497-kxl46
                                 1/1
                                         Running
                                                                2m10s
                                                                        10.40.0.1
                                                                                     k8s-worker1
                                                                                                    <none>
                                                                                                                      <none
guru@k8s-master:~$ kubectl get no
              STATUS
                                           ROLES
                                                            AGE
                                                                    VERSION
NAME
k8s-master
                                           control-plane
                                                            293d
              Ready
                                                                    v1.28.8
              Ready SchedulingDisabled
k8s-worker1
                                           <none>
                                                            293d
                                                                    v1.28.8
k8s-worker2
                                                            293d
                                                                    v1.28.8
```

## 노드관리

## 6. Pod Scheduling

다음의 조건으로 pod를 생성하세요.

- Name: eshop-store
- Image: nginx
- Nodeselector: disktype=ssd

worker1,2 노드에 라벨링 추가

kubectl label node worker1 disktype=ssd

kubectl label node worker2 disktype=hdd

```
guru@k8s-master:~$ kubectl label node k8s-worker1 disktype=ssd
node/k8s-worker1 labeled
guru@k8s-master:~$ kubectl label node k8s-worker2 disktype=hdd
node/k8s-worker2 labeled
```

## 라벨링 추가 여부 확인

```
guru@k8s-master:~$ kubectl get no -L disktype
                                                           DISKTYPE
NAME
               STATUS
                        ROLES
                                         AGE
                                                 VERSION
                                                 v1.28.8
k8s-master
              Ready
                        control-plane
                                         293d
k8s-worker1
                                         293d
                                                 v1.28.8
              Ready
                        <none>
                                                           ssd
                                                 v1.28.8
                                                           hdd
k8s-worker2
              Ready
                                         293d
                        <none>
```

pod 생성을 위한 yaml파일 작성

```
guru@k8s-master:~$ kubectl run eshop-store --image=nginx --dry-run=client -o yaml > eshop-store.yaml guru@k8s-master:~$
```

vi eshop-store.yaml 수정

## 수정 전

```
apiVersion: v1
kind: Pod
metadata:
    creationTimestamp: null
    labels:
        run: eshop-store
    name: eshop-store
spec:
    containers:
    - image: nginx
        name: eshop-store
        resources: {}
    dnsPolicy: ClusterFirst
    restartPolicy: Always
status: {}
```

## 수정 후

```
apiVersion: v1
kind: Pod
metadata:
   name: eshop-store
spec:
   containers:
   - image: nginx
     name: eshop-store
nodeSelector:
     disktype: ssd
```

yaml 적용

kubectl apply -f eshop-store.yaml

```
guru@k8s-master:~$ kubectl apply -f eshop-store.yaml
pod/eshop-store created
```

pod 확인

kubectl get po eshop-store -o wide

```
guru@k8s-master:~$ kubectl get po eshop-store
NAME
              READY
                                                  ΙP
                                                               NODE
                                                                             NOMINATED NODE
                                                                                               READINESS G
                      STATUS
                                 RESTARTS
                                            AGE
ATES
              1/1
                      Running
                                 0
                                            70s
                                                  10.40.0.1 k8s-worker1
eshop-store
                                                                             <none>
                                                                                               <none>
```

# 파드 생성

# 7. 환경변수, command, args 적용

'cka-exam'이라는 namespace를만들고, 'cka-exam' namespace에 아래와 같은 Pod를 생성하시오.

- pod Name: pod-01
- image: busybox
- 환경변수 : CERT = "CKA-cert"
- ommand: /bin/sh
- args: "-c", "while true; do echo \$(CERT); sleep 10;done"

네임스페이스 생성

kubectl create ns cka-exam

```
guru@k8s-master:~$ kubectl create ns cka-exam
namespace/cka-exam created
```

## 네임스페이스 확인

```
guru@k8s-master:~$ kubectl get ns
NAME
                   STATUS
                             AGE
cka-exam
                             105s
                   Active
default
                             293d
                   Active
kube-node-lease
                             293d
                   Active
kube-public
                   Active
                             293d
                             293d
kube-system
                   Active
```

```
guru@k8s-master:~$ kubectl run pod-01 --image=busybox -n cka-exam --env=CERT=CKA-cert --dry-run=client -o yaml > pod -01.yaml guru@k8s-master:~$ [
```

vi pod-01.yaml 확인 및 수정

- pod Name: pod-01

- image: busybox

- 환경변수 : CERT = "CKA-cert"

- ommand: /bin/sh

- args: "-c", "while true; do echo \$(CERT); sleep 10;done"

## 수정 전

## apiVersion: vl kind: Pod metadata: creationTimestamp: null labels: run: pod-01 name: pod-01 namespace: cka-exam spec: containers: - env: - name: CERT value: CKA-cert image: busybox name: pod-01 resources: {} dnsPolicy: ClusterFirst restartPolicy: Always status: {}

수정 후

```
apiVersion: v1
kind: Pod
metadata:
   name: pod-01
namespace: cka-exam
spec:
   containers:
   - env:
        - name: CERT
        value: CKA-cert
        command: [/bin/sh]
        args: ["-c", "while true; do echo $(CERT); sleep 10;done"]
        image: busybox
        name: pod-01
```

yaml 적용

kubectl apply -f pod-01.yaml

```
guru@k8s-master:~$ kubectl apply -f pod-01.yaml
pod/pod-01 created
```

pod 확인

kubectl get pod-01 -n cka-exam -o wide

```
guru@k8s-master:~$ kubectl get po pod-01 -n cka-exam -o
                                                          wide
NAME
         READY
                  STATUS
                            RESTARTS
                                       AGE
                                                           NODE
                                                                         NOMINATED NODE
                                                                                           READINESS GATES
pod-01
                                       113s
                                              10.32.0.2
                                                           k8s-worker2
         1/1
                 Running
                            0
                                                                         <none>
                                                                                           <none>
```

## 파드 생성

## 8. Static Pod 생성

worker1 노드에 nginx-static-pod.yaml 라는 이름의 Static Pod를 생성하세요.

pod name : nginx-static-pod

- image : nginx

- port: 80

## worker 1 접속

```
guru@k8s-master:~$ ssh k8s-worker1
The authenticity of host 'k8s-workerl (10.100.0.106)' can't be established.
ECDSA key fingerprint is SHA256:p6Kg7OSM9ozOHdkSA5fz9qlEScr9wis1+l7XjJpNtCI.
Are you sure you want to continue connecting (yes/no/[fingerprint])? yes
Warning: Permanently added 'k8s-worker1,10.100.0.106' (ECDSA) to the list of known hosts.
guru@k8s-worker1's password:
Welcome to Ubuntu 20.04.6 LTS (GNU/Linux 5.15.0-79-generic x86 64)
 * Documentation: https://help.ubuntu.com
 * Management:
                   https://landscape.canonical.com
 * Support:
                   https://ubuntu.com/advantage

    Introducing Expanded Security Maintenance for Applications.

   Receive updates to over 25,000 software packages with your
   Ubuntu Pro subscription. Free for personal use.
     https://ubuntu.com/pro
Expanded Security Maintenance for Applications is not enabled.
219 updates can be applied immediately.
174 of these updates are standard security updates.
추가 업데이트를 확인하려면 apt list --upgradable 을 실행하세요.
Enable ESM Apps to receive additional future security updates.
See https://ubuntu.com/esm or run: sudo pro status
The list of available updates is more than a week old.
To check for new updates run: sudo apt update
Your Hardware Enablement Stack (HWE) is supported until April 2025.
Last login: Mon Apr _1 11:03:01 2024 from 10.100.0.2
guru@k8s-worker1:~$
```

worker1의 static-pod 위치 확인

- 1. cd /var/lib/kubelet 이동
- 2. cat config.yaml | grep -i static 으로 위치 확인
- 3. cd /etc/kubernetes/manifests 이동

```
guru@k8s-worker1:~$ cd /var/lib/kubelet
guru@k8s-worker1:/var/lib/kubelet$ cat config.yaml | grep -i static
staticPodPath: /etc/kubernetes/manifests
guru@k8s-worker1:/var/lib/kubelet$ cd /etc/kubernetes/manifests
guru@k8s-worker1:/etc/kubernetes/manifests$
```

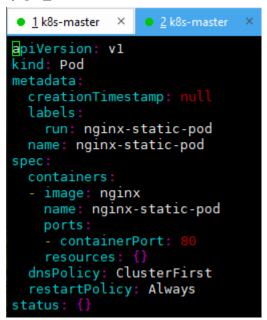
(추가 세션) master 창에서 nginx-static-pod.yaml 파일 생성 kubectl run nginx-static-pod --image=nginx --port=80 --dry-run=client -o yaml > nginx-static-pod.yaml

```
• 1 k8s-master × • 2 k8s-master × +

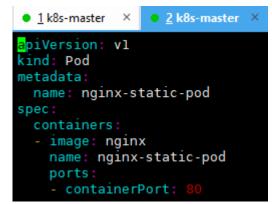
guru@k8s-master:~$ kubectl run nginx-static-pod --image=nginx --port=80 --dry-run=client -o yaml > nginx-static-pod.
yaml
guru@k8s-master:~$ [
```

vi nginx-static-pod.yaml 수정 - worker1 노드에 수정된 코드 내용을 붙여넣을 예정

## 수정 전







- pod name : nginx-static-pod

- image : nginx

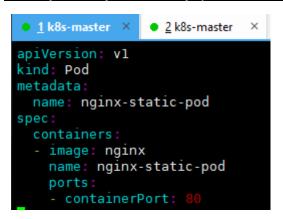
- port: 80

worker1 에서 yaml 파일 생성

[\*\* worker1 사용자 계정으로는 입력한 내용이 저장되지 않아 root 계정으로 진입 \*\*]

sudo vi nginx-static-pod.yaml (수정한 코드 붙여넣기)

## guru@k8s-worker1:/etc/kubernetes/manifests\$ sudo vi nginx-static-pod.yaml [sudo] guru 암호:



- (추가세션) master 창에서 수정된 코드 복사 후 worker1에서 실행한 vi nginx-static-pod. yaml 안에 코드 붙여넣기
- 이 후 마스터 창 vi nginx-static-pod. yaml 나가기 :q!

master 노드에서 'nginx-static-pod-k8s-worker1' pod 확인

kubectl get po nginx-static-pod-k8s-worker1 -o wide

```
guru@k8s-master:~$ kubectl get po nginx-static-pod-k8s-workerl -o wide
NAME
                                READY
                                        STATUS
                                                   RESTARTS
                                                                   AGE
                                                                                      NODE
                                                                                                    NOMINATED NODE
                                                                                                                      REA
DINESS GATES
nginx-static-pod-k8s-workerl
                                                                                     k8s-worker1
                                1/1
                                        Running
                                                   2 (9m9s ago)
                                                                         10.40.0.2
                                                                                                    <none>
                                                                                                                      <nc
```

## 파드 생성

## 9. 로그 확인

Pod "nginx-static-pod-k8s-worker1"의 log를 모니터링하고, 메세지를 포함하는 로그라인을 추출하세요. 추출된 결과는 /opt/REPORT/2023/pod-log에 기록하세요.

nginx-static-pod-worker1 작동 여부 확인

kubectl get po nginx-static-pod-worker1

```
guru@k8s-master:~$ kubectl get po nginx-static-pod-k8s-workerl
NAME READY STATUS RESTARTS AGE
nginx-static-pod-k8s-workerl 1/1 Running 3 (2m57s ago) 37m
```

root 계정 전환 (sudo -i)

/opt/REPORT/2023 디렉터리 생성

```
guru@k8s-master:~$ sudo -i
[sudo] guru 앞 호 :
root@k8s-master:~# mkdir -p /opt/REPORT/2023/
root@k8s-master:~#
```

#### worker1 log 모니터링

```
root@k8s-master:~# mkdir -p /opt/REPORT/2023/
root@k8s-master:~# kubectl logs nginx-static-pod-k8s-workerl
/docker-entrypoint.sh: /docker-entrypoint.d/ is not empty, will attempt to perform configuration
/docker-entrypoint.sh: Looking for shell scripts in /docker-entrypoint.d/
/docker-entrypoint.sh: Launching /docker-entrypoint.d/10-listen-on-ipv6-by-default.sh
10-listen-on-ipv6-by-default.sh: info: Getting the checksum of /etc/nginx/conf.d/default.conf
10-listen-on-ipv6-by-default.sh: info: Enabled listen on IPv6 in /etc/nginx/conf.d/default.conf
docker-entrypoint.sh: Sourcing /docker-entrypoint.d/15-local-resolvers.envsh/
/docker-entrypoint.sh: Launching /docker-entrypoint.d/20-envsubst-on-templates.sh
/docker-entrypoint.sh: Launching /docker-entrypoint.d/30-tune-worker-processes.sh
/docker-entrypoint.sh: Configuration complete; ready for start up
2025/01/19 15:44:13 [notice] 1#1: using the "epoll" event method
2025/01/19 15:44:13 [notice] 1#1: nginx/1.27.3
2025/01/19 15:44:13 [notice] 1#1: built by gcc 12.2.0 (Debian 12.2.0-14)
2025/01/19 15:44:13 [notice] 1#1: OS: Linux 5.15.0-79-generic
2025/01/19 15:44:13 [notice] 1#1: getrlimit(RLIMIT_NOFILE): 1048576:1048576
2025/01/19 15:44:13 [notice] 1#1: start worker processes 2025/01/19 15:44:13 [notice] 1#1: start worker process 29
2025/01/19 15:44:13 [notice] 1#1: start worker process 30
root@k8s-master:~#
```

kubectl logs nginx-static-pod-worker1 > /opt/REPORT/2023/pod-log

cat /opt/REPORT/2023/pod-log

```
root@k8s-master:~# kubectl logs nginx-static-pod-k8s-worker1 > /opt/REPORT/2023/pod-log
root@k8s-master:~#
root@k8s-master:~# cat /opt/REPORT/2023/pod-log
/docker-entrypoint.sh: /docker-entrypoint.d/ is not empty, will attempt to perform configuration
/docker-entrypoint.sh: Looking for shell scripts in /docker-entrypoint.d/
/docker-entrypoint.sh: Launching /docker-entrypoint.d/10-listen-on-ipv6-by-default.sh
10-listen-on-ipv6-by-default.sh: info: Getting the checksum of /etc/nginx/conf.d/default.conf 10-listen-on-ipv6-by-default.sh: info: Enabled listen on IPv6 in /etc/nginx/conf.d/default.conf
/docker-entrypoint.sh: Sourcing /docker-entrypoint.d/15-local-resolvers.envsh
/docker-entrypoint.sh: Launching /docker-entrypoint.d/20-envsubst-on-templates.sh
/docker-entrypoint.sh: Launching /docker-entrypoint.d/30-tune-worker-processes.sh
/docker-entrypoint.sh: Configuration complete; ready for start up
2025/01/19 15:44:13 [notice] 1#1: using the "epoll" event method
2025/01/19 15:44:13 [notice] 1#1: nginx/1.27.3
2025/01/19 15:44:13 [notice] 1#1: built by gcc 12.2.0 (Debian 12.2.0-14)
2025/01/19 15:44:13 [notice] 1#1: 0S: Linux 5.15.0-79-generic
2025/01/19 15:44:13 [notice] 1#1: getrlimit(RLIMIT NOFILE): 1048576:1048576
2025/01/19 15:44:13 [notice] 1#1: start worker processes
2025/01/19 15:44:13 [notice] 1#1: start worker process 29
2025/01/19 15:44:13 [notice] 1#1: start worker process 30
```

# 파드 생성

## 10. 로그 확인

4개의 컨테이너를 동작시키는 eshop-frontend Pod를 생성하시오.

- pod image: nginx, redis, memcached, consul

eshop-frontend Pod 생성 및 적용 / 확인

kubectl run eshop-frontend --image=nginx --dry-run=client -o yaml > eshop-frontend.yaml vi eshop-frontend.yaml

```
apiVersion: v1
kind: Pod
metadata:
   name: eshop-frontend
spec:
   containers:
   - image: nginx
   name: nginx
   - image: redis
   name: redis
   - image: memcached
   name: memcached
   image: consul
   name: consul
```

kubectl apply -f eshop-frontend.yaml kubectl get po pod image:

- nginx
- redis
- memcached
- consul

```
guru@k8s-master:~$ kubectl run eshop-frontend --image=nginx --dry-run=client -o yaml > eshop-frontend.yaml
guru@k8s-master:~$ vi eshop-frontend.yaml
guru@k8s-master:~$ kubectl apply -f eshop-frontend.yaml
pod/eshop-frontend created
guru@k8s-master:~$ kubectl get po
NAME
                                READY
                                        STATUS
                                                            RESTARTS
                                                                           AGE
eshop-frontend
                                        ContainerCreating
                                0/4
                                                                           11s
nginx-static-pod-k8s-worker1
                                                            3 (10m ago)
                                                                           3h40m
                                1/1
                                        Running
guru@k8s-master:~$ kubectl get po
                                                           RESTARTS
NAME
                                READY
                                        STATUS
                                                                          AGE
                                        ImagePullBackOff
eshop-frontend
                                                                          108s
                                3/4
                               1/1
nginx-static-pod-k8s-worker1
                                                                          3h42m
                                        Running
                                                           3 (12m ago)
```

# 쿠버네티스 컨트롤러와 네트워킹

## 11. Rolling Update & Roll Back

- 1. Deployment를 이용해 nginx 파드를 3개 배포한 다음 컨테이너 이미지 버전을 rolling update하고 update record를 기록합니다.
- 2. 마지막으로 컨테이너 이미지를 previous version으로 roll back 합니다.
  - name: eshop-payment
  - Image : nginx
  - Image version: 1.16
  - update image version: 1.17
  - label: app=payment, environment=production
- 1. Deployment를 이용해 nginx 파드를 3개 배포한 다음 컨테이너 이미지 버전을 rolling update하고 update record를 기록

kubectl create deployment eshop-payment --image=nginx:1.16 --replicas=3 --dry-run=client -o yaml > eshop-payment.yaml

guru@k8s-master:~\$ kubectl create deployment eshop-payment --image=nginx:1.16 --replicas=3 --dry-run=client -o yaml > eshop-payment.yaml

vi eshop eshop-payment.yaml

```
apiVersion: apps/v1
                                                 - name: eshop-payment
kind: Deployment
                                                 - Image: nginx
metadata:
                                                 - Image version: 1.16
  labels:
    app: payment
                                                 - update image version: 1.17
    environment: production
                                                 - label: app=payment, environment=production
  name: eshop-payment
spec:
  replicas: 3
  selector:
    matchLabels:
      app: payment
      environment: production
  template:
    metadata:
      labels:
         app: payment
         environment: production
    spec:
      containers:
       - image: nginx:1.16
         name: nginx
```

kubectl apply -f eshop-payment.yaml -record

```
guru@k8s-master:~$ kubectl apply -f eshop-payment.yaml --record Flag --record has been deprecated, --record will be removed in the future deployment.apps/eshop-payment created
```

kubectl set image deploy eshop-payment nginx=nginx:1.17 -record

```
guru@k8s-master:~$ kubectl set image deploy eshop-payment nginx=nginx:1.17 --record Flag --record has been deprecated, --record will be removed in the future deployment.apps/eshop-payment image updated
```

버전 확인

kubectl get deploy,po | grep -i eshop-payment kubectl describe po eshop-payment-5c4cc77dcd-58fb2 | grep -i nginx kubectl rollout history deploy eshop-payment

```
uru@k8s-master:~$ kubectl get
deployment.apps/eshop-payment 3/3
pod/eshop-payment-5c4cc77dcd-58fb2
pod/eshop-payment-5c4cc77dcd-gj46h
pod/eshop-payment-5c4cc77dcd-rd8kp
                                       3/3
                                              1/1
                                                       Running
                                                                   Θ
                                                                                     20s
                                             1/1
                                                       Running
                                                                   0
                                                                                     28s
                                                       Running
                                             1/1
                                                                   Θ
                                                                                     17s
guru@k8s-master:~$ kubectl describe po eshop-payment-5c4cc77dcd-58fb2 | grep -i nginx
     Image:
                         nginx:1.17
                         docker.io/library/nginx@sha256:6fff55753e3b34e36e24e37039ee9eae1fe38a6420d8ae16ef37c92d1eb26699
     Image ID:
                                                        Container image "nginx:1.17" already present on machine Created container nginx
  Normal Pulled
                                 kubelet
                         70s
  Normal
            Created
                          70s
                                 kubelet
  Normal
                                 kubelet
                                                         Started container nginx
           Started
                          70s
guru@k8s-master:~$ kubectl rollout history deploy eshop-payment
deployment.apps/eshop-payment
REVISION CHANGE-CAUSE
            kubectl apply --filename=eshop-payment.yaml --record=true
            kubectl set image deploy eshop-payment nginx=nginx:1.17 --record=true
```

2. 컨테이너 이미지를 previous version으로 roll back

kubectl rollout undo deploy eshop-payment

guru@k8s-master:~\$ kubectl rollout undo deploy eshop-payment deployment.apps/eshop-payment rolled back

확인

kubectl get po

kubectl describe po eshop-payment-7ddf6d4467-ddxsk | grep -i nginx

```
guru@k8s-master:~$ kubectl get po
                                    READY
                                                       RESTARTS
NAME
                                             STATUS
                                                                       AGE
eshop-payment-7ddf6d4467-ddxsk
                                             Running
                                                       Θ
                                                                       36s
eshop-payment-7ddf6d4467-pfwnq
                                             Running
                                                       0
                                                                       35s
eshop-payment-7ddf6d4467-v5w5c
                                    1/1
                                             Running
                                                       0
                                                                       33s
                                    1/1
                                            Running
                                                       3 (70m ago)
nginx-static-pod-k8s-workerl
                                                                       11m
guru@k8s-master:~$ kubectl describe po eshop-payment-7ddf6d4467-ddxsk | grep -i nginx
    Image:
                      docker.io/library/nginx@sha256:d20aa6d1cae56fd17cd458f4807e0de462caf2336f0b70b5eeb69fcaaf30dd9c
    Image ID:
                                                  Container image "nginx:1.16" already present on machine Created container nginx
                             kubelet
kubelet
          Pulled
  Normal
                       74s
                       74s
  Normal
          Created
                                                  Started container nginx
                             kubelet
```

# 쿠버네티스 컨트롤러와 네트워킹

## 12. ClusterIP

- 1. 'devops' namespace에서 deployment eshop-order를 다음 조건으로 생성하시오.
  - image: nginx, replicas: 2, label: name=order
- 2. 'eshop-order' deployment의 Service를 만드세요.
  - Service Name: eshop-order-svc
  - Type: ClusterIP
  - Port: 80

kubectl get ns

kubectl create ns devops

kubectl get ns

kubectl create deploy eshop-order -n devops --image=nginx --replicas=2 --dry-run=client -o yaml > eshop-order.yaml

vi eshop-order.yaml

kubectl apply -f eshop-order.yaml

kubectl expose deploy eshop-order -n devops --name=eshop-order-svc --port=80 --target-port=80

kubectl get deploy eshop-order -n devops

kubectl get svc eshop-order-svc -n devops

```
guru@k8s-master:~$ kubectl get
NAME
                  STATUS
                           AGE
default
                  Active
                           290d
kube-node-lease
                 Active
                           290d
kube-public
                 Active
                           290d
kube-system
                 Active
                           290d
guru@k8s-master:~$ kubectl create ns devops
namespace/devops created
guru@k8s-master:~$ kubectl get ns
NAME
                 STATUS AGE
default
                 Active
                           290d
devops
                 Active
                           7s
kube-node-lease Active
                           290d
kube-public
                 Active
                           290d
                 Active
                           290d
kube-system
guru@k8s-master:~$ kubectl create deploy eshop-order -n devops --image=nginx --replicas=2 --dry-run=client -
o yaml > eshop-order.yaml
guru@k8s-master:~$ kubectl apply -f eshop-order.yaml
deployment.apps/eshop-order created
guru@k8s-master:~$ vi eshop-order.yaml
guru@k8s-master:~$ kubectl apply -f eshop-order.yaml
deployment.apps/order created
guru@k8s-master:~$ kubectl espose deploy eshop-order -n devops --name=eshop-order-svc --port=80 --target-por
t=80
error: unknown command "espose" for "kubectl"
Did you mean this?
        expose
guru@k8s-master:~$ kubectl expose deploy eshop-order -n devops --name=eshop-order-svc --port=80 --target-por
t=80
service/eshop-order-svc exposed
guru@k8s-master:~$ kubectl get deploy,svc eshop-order -n devops
NAME
              READY
                     UP-TO-DATE
                                   AVAILABLE
                                               AGE
              2/2
                                               5m38s
eshop-order
Error from server (NotFound): services "eshop-order" not found
guru@k8s-master:~$ kubectl get deploy eshop-order -n devops
              READY
                      UP-TO-DATE
                                   AVAILABLE
                                               AGE
                                               5m52s
eshop-order
              2/2
                      2
                                   2
guru@k8s-master:~$ kubectl get svc eshop-order-svc -n devops
                              CLUSTER-IP
                                              EXTERNAL-IP
                                                            PORT(S)
                                                                      AGE
                  ClusterIP
                              10.98.103.237
                                                            80/TCP
                                                                      59s
eshop-order-svc
                                              <none>
```

# 쿠버네티스 컨트롤러와 네트워킹

# 13. NodePort

- 1. 'front-end' deployment를 다음 조건으로 생성하시오.
  - image: nginx, replicas: 2, label: run=nginx
- 2. 'front-end' deployment의 nginx 컨테이너를 expose하는 'front-end-nodesvc'라는 새 service를 만듭니다.
- 3. Front-end로 동작중인 Pod에는 node의 \*\*30200\*\* 포트로 접속되어야 합니다.

kubectl create deploy front-end --image=nginx --replicas=2 --dry-run=client -o yaml > front-end,yaml

vi front-end,yaml

kubectl expose deploy front-end --name= front-end-nodesvc --type=NodePort --port=80 --target-port=80 --dry-run=client -o yaml > front-end-nodesvc.yaml

vi front-end-nodesvc.yaml

kubectl apply -f vi front-end-nodesvc.yaml

kubectl get deploy,svc

vi front-end-nodesvc.yaml

kubectl apply -f front-end-nodesvc.yaml

kubectl get deploy,svc

# **NetworkPolicy & Ingress**

## 14. Network Policy (정책)

- 1. customera, customerb를 생성한 후, 각각 PARTITION=customera, PARTITION=customerb를 라벨링하시오.
- 2. default namespace에 다음과 같은 pod를 생성하세요.
- name: pocimage: nginxport: 80
- label: app=poc
- "partition=customera"를 사용하는 namespace에서만 poc의
- 3. 80포트로 연결할 수 있도록 default namespace에 'allow-web-from-customera'라는 network Policy를 설정하세요.보안 정책상 다른 namespace의 접근은 제한합니다.

## ☑ customera, customerb를 생성한 후,

각각 PARTITION=customera, PARTITION=customerb를 라벨링하시오.

✔ customera, customerb 네임스페이스 생성

```
guru@k8s-master:~$ kubectl create ns customera namespace/customera created guru@k8s-master:~$ kubectl create ns customerb namespace/customerb created guru@k8s-master:~$
```

## ✔ 네임스페이스 확인

```
guru@k8s-master:~$ kubectl get ns
NAME
                  STATUS
                           AGE
api-access
                  Active
                           3d20h
                  Active
                           2m31s
customera
customerb
                  Active
                           2m29s
default
                  Active
                           291d
devops
                  Active
                           2d20h
ing-internal
                           19h
                  Active
kube-node-lease
                  Active
                           291d
kube-public
                  Active
                           291d
                           291d
kube-system
                  Active
guru@k8s-master:~$
```

✔ customera, customerb 파티션 및 라벨링

```
guru@k8s-master:~$ kubectl label ns customera partition=customera
namespace/customera labeled
guru@k8s-master:~$
guru@k8s-master:~$ kubectl label ns customerb partition=customerb
namespace/customerb labeled
```

## ✔ 네임스페이스 Partition 및 라벨 확인

```
guru@k8s-master:~$ kubectl get ns
NAME
                   STATUS
                             AGE
                             3d20h
api-access
                   Active
                   Active
                             20s
customera
                             19s
customerb
                   Active
default
                   Active
                             291d
devops
                   Active
                             2d20h
ing-internal
                   Active
                             19h
kube-node-lease
                   Active
                             291d
kube-public
                   Active
                             291d
kube-system
                             291d
                   Active
guru@k8s-master:~$ kubectl get ns -L partition
NAME
                   STATUS
                             AGE
                                     PARTITION
api-access
                             3d20h
                   Active
                             25s
customera
                   Active
                                     customera
customerb
                   Active
                             24s
                                     customerb
default
                   Active
                             291d
devops
                   Active
                             2d20h
ing-internal
                   Active
                             19h
kube-node-lease
                   Active
                             291d
kube-public
                   Active
                             291d
kube-system
                             291d
                   Active
guru@k8s-master:~$
```

Partiton 생성과 라벨 customera, customer가 생성된 것을 확인할 수 있다.

# 📝 default namespace에 다음과 같은 pod를 생성하세요

❤ poc 생성

guru@k8s-master:~\$ kubectl run poc --image=nginx --port=80 --labels=app=poc
pod/poc created

✓ pod 확인

guru@k8s-master:~\$ kubectl get po poc NAME READY STATUS RESTARTS AGE poc 0/1 Pending 0 5s

✔ Kubernetes.io/docs 내 Network Policy 검색



Q Search this site

About 4,850 results (0.11 seconds)

## Network Policies | Kubernetes

Kubernetes > docs > concepts > services-networking > network-policies

Apr 1, 2024 ... NetworkPolicies are an application-centric construct which allow you to specif

## ✔ Network Policy Recource 검색

```
service/networking/networkpolicy.yaml
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: test-network-policy
 namespace: default
spec:
  podSelector:
   matchLabels:
      role: db
  policyTypes:
  - Ingress
  - Egress
  ingress:
  - from:
    - ipBlock:
        cidr: 172.17.0.0/16
        except:
        - 172.17.1.0/24
    - namespaceSelector:
        matchLabels:
          project: myproject
    - podSelector:
        matchLabels:
          role: frontend
    ports:
    - protocol: TCP
      port: 6379
  egress:
  - to:
    - ipBlock:
        cidr: 10.0.0.0/24
    ports:
    - protocol: TCP
      port: 5978
```

☑ 80포트로 연결할 수 있도록 default namespace에 'allow-web-from-customera'라는 network Policy를 설정하세요. 보안 정책 상 다른 namespace의 접근은 제한합니다.

✔ vi netpol.yaml 생성 및 수정

```
수정 전
● 1 k8s-master × ● 2 k8s-master × +
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: test-network-policy
  namespace: default
spec:
  podSelector:
  matchLabels:
     role: db
  policyTypes:
  - Ingress

    Egress

  ingress:
  - from:
    ipBlock:
        cidr: 172.17.0.0/16
        except:
        - 172.17.1.0/24

    namespaceSelector:

        matchLabels:
          project: myproject
    podSelector:
        matchLabels:
          role: frontend
    ports:
    - protocol: TCP
      port: 6379
  egress:
  - to:
    ipBlock:
        cidr: 10.0.0.0/24
    ports:

    protocol: TCP

      port: 5978
```

```
수정 후

    1 k8s-master ×  ■ 2 k8s-master × +

apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: allow-web-from-customera
  namespace: default
spec:
  podSelector:
    matchLabels:
      app: poc
  policyTypes:

    Ingress

  ingress:
  - from:

    namespaceSelector:

        matchLabels:
           partition: customera
    ports:
    - protocol: TCP
      port: 80
```

✓ yaml 적용

```
guru@k8s-master:~$ kubectl apply -f netpol.yaml
networkpolicy.networking.k8s.io/allow-web-from-customera unchanged
```

✓ pod 확인

```
guru@k8s-master:~$ kubectl get netpol
NAME POD-SELECTOR AGE
allow-web-from-customera app=poc 47h
```

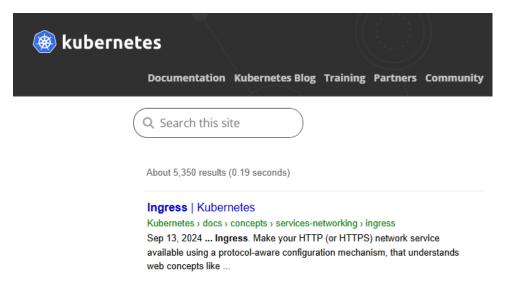
# 15. Ingress

- 1. Create a new nginx Ingress resource as follows
  - Name: ping
  - Namespace : ing-internal
  - Exposing service hi on path /hi using service port 5678



## **Treate a new nginx Ingress resource as follows**

✔ Kubenates 내 ingress 검색



# The Ingress resource

A minimal Ingress resource example:

```
service/networking/minimal-ingress.yaml
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: minimal-ingress
  annotations:
    nginx.ingress.kubernetes.io/rewrite-target: /
  {\bf ingressClassName} \colon \; {\tt nginx-example} \\
  rules:
  - http:
      paths:
      - path: /testpath
        pathType: Prefix
        backend:
           service:
             name: test
             port:
               number: 80
```

✔ 네임스페이스 생성

guru@k8s-master:~\$ kubectl create ns ing-internal namespace/ing-internal created

## ✔ 네임스페이스 확인

```
guru@k8s-master:~$ kubectl get ns ing-internal
NAME STATUS AGE
ing-internal Active 29s
```

## ✓ vi ingress.yaml

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: ping
  namespace: ing-internal
  ingressClassName: nginx-example
  rules:
  - http:
      paths:
      - path: /hi
        pathType: Prefix
        backend:
          service:
            name: hi
            port:
              number: 5678
```

Name : ping

Namespace : ing-internal

 Exposing service hi on path /hi using service port 5678

- Service name : hi

- Path : /hi- port : 5679

apply -f 적용 내용포함 요망

## ✔ Ingress 확인

```
guru@k8s-master:~$ kubectl get Ingress -n ing-internal NAME CLASS HOSTS ADDRESS PORTS AGE ping nginx-example * 80 5m18s
```

## 16. Service and DNS Lookup

- 1. image nginx를 사용하는 resolver pod를 생성하고 resolver-service라는 service를 구성합니다.
- 2. 클러스터 내에서 service와 pod 이름을 조회할 수 있는지 테스트합니다.
  - dns 조회에 사용하는 pod 이미지는 busybox:1.28이고, service와 pod 이름 조회는 nlsookup을 사용합니다.
  - service 조회 결과는 /var/CKA2023/nginx.svc에 pod name 조회 결과는 /var/CKA2023/nginx.pod 파일에 기록합니다.

## 

- 📝 image nginx를 사용하는 resolver pod를 생성하고 resolver-service라는 service를 구성합니다.
- ✔ Resolver pod를 생성

guru@k8s-master:~\$ kubectl run resolver --image=nginx --port=80
pod/resolver created

✔ pod 확인

guru@k8s-master:~\$ kubectl get po NAME READY STATUS RESTARTS AGE resolver 1/1 Running 0 60s

✔ 서비스 생성 kubectl expose pod resolver --name=resolver-service --port=80

guru@k8s-master:~\$ kubectl expose pod resolver --name=resolver-service --port=80
service/resolver-service exposed

✔ 서비스 확인 kubectl get svc resolver-service

guru@k8s-master:~\$ kubectl get svc resolver-service
NAME TYPE CLUSTER-IP EXTERNAL-IP PORT(S) AGE
resolver-service ClusterIP 10.100.54.140 <none> 80/TCP 84s

# ☑ 클러스터 내에서 service와 pod 이름을 조회할 수 있는지 테스트합니다.

✓ Run test-nslookup

```
guru@k8s-master:~$ sudo -i
root@k8s-master:~# mkdir -p /var/CKA2023
root@k8s-master:~# kubectl run test-nslookup --image=busybox:1.28 -it --rm --restart=Never -- nslookup 10.100.54.140
Server: 10.96.0.10
Address 1: 10.96.0.10 kube-dns.kube-system.svc.cluster.local

Name: 10.100.54.140
Address 1: 10.100.54.140 resolver-service.default.svc.cluster.local
pod "test-nslookup" deleted
root@k8s-master:~# kubectl run test-nslookup --image=busybox:1.28 -it --rm --restart=Never -- nslookup 10.100.54.140 > /var/CKA2023/nginx.svc
```

root 전환 및 mkdir -p

```
guru@k8s-master:~$ sudo -i
root@k8s-master:~# mkdir -p /var/CKA2023
```

run test-nslookup 10.100.54.140

```
root@k8s-master:~# kubectl run test-nslookup --image=busybox:1.28 -it --rm --restart=Never -- nslookup 10.100.54.140
Server: 10.96.0.10
Address 1: 10.96.0.10 kube-dns.kube-system.svc.cluster.local

Name: 10.100.54.140
Address 1: 10.100.54.140 resolver-service.default.svc.cluster.local
pod "test-nslookup" deleted
```

run test-nslookup 10.100.54.140 -image=busybox:1.28 -it --rm --restart=Never - resolver 10.100.54.140 > nginx.svc

```
root@k8s-master:~# kubectl run test-nslookup --image=busybox:1.28 -it --rm --restart=Never -- nslookup 10.100.54.140 > nginx.svc root@k8s-master:~#
```

cat /var/cka2023/nginx.svc

```
root@k8s-master:~# cat /var/CKA2023/nginx.svc
Server: 10.96.0.10
Address 1: 10.96.0.10 kube-dns.kube-system.svc.cluster.local

Name: 10.100.54.140
Address 1: 10.100.54.140 resolver-service.default.svc.cluster.local
pod "test-nslookup" deleted
```

get po resolver -o wide

```
get po resolver
                    kubectl
                                                   ΤP
NAME
            READY
                    STATUS
                               RESTARTS
                                           AGE
                                                                NODE
                                                                               NOMINATED NODE
                                                                                                 READINESS GATES
                                                   10.46.0.2
                                           3m54s
resolver
            1/1
                    Running
                               0
                                                                k8s-worker1
                                                                               <none>
                                                                                                 <none>
```

kube run test-nslookup –image=busybox:1.28 -it --rm --restart=Never – nslookp 10-46-0-2 default.pod.cluster.local kube run test-nslookup –image=busybox:1.28 -it --rm --restart=Never – nslookp 10-46-0-2 default.pod.cluster.local > /var/CKA2023/nginx.pod

```
root@k8s-master:-# kubectl run test-nslookup --image=busybox:1.28 -it --rm --restart=Never -- nslookup 10-46-0-2.default.pod.cluster.local
If you don't see a command prompt, try pressing enter.
warning: couldn't attach to pod/test-nslookup, falling back to streaming logs: Internal error occurred: error attaching to container: failed to load task: no runninc
Server: 10.96.0.10
Redress 1: 10.96.0.10
Redress 1: 10.96.0.10 kube-dns.kube-system.svc.cluster.local
Redress 1: 10.46-0-2.default.pod.cluster.local
Redress 1: 10.46.0.2 10-46-0-2.resolver-service.default.svc.cluster.local
```

root@k8s-master:~# cat /var/CKA2023/nginx.pod

Server: 10.96.0.10

Address 1: 10.96.0.10 kube-dns.kube-system.svc.cluster.local

Name: 10-46-0-2.default.pod.cluster.local

Address 1: 10.46.0.2 10-46-0-2.resolver-service.default.svc.cluster.local

pod "test-nslookup" deleted

# Resoure 관리

# 17. emptyDir Volume

1. 다음 조건에 맞춰서 nginx 웹서버 pod가 생성한 로그파일을 받아서 STDOUT으로 출력하는 busybox 컨테이너를 운영하시오.

Pod Name: \*\*weblog\*\*

- \*\*Web container:\*\*
- Image: \*\*nginx:1.17\*\*
- Volume mount: \*\*/var/log/nginx\*\*
- Readwrite

Log container:

- Image: busybox
- args: /bin/sh, -c, "tail -n+1 -f /data/access.log"
- Volume mount : /data
- readonly
- 2. emptyDir 볼륨을 통한 데이터 공유

☑ 다음 조건에 맞춰서 nginx 웹서버 pod가 생성한 로그파일을 받아서 STDOUT으로 출력하는 busybox 컨테이너를 운영하시오.

쿠버네티스 검색(emptyDir)



**Documentation** 

Q Search this site

About 5,160 results (0.15 seconds)

## Volumes | Kubernetes

Kubernetes > docs > concepts > storage > volumes

Dec 18, 2024 ... Kubernetes volumes provide a way for containers in a pods to access and s

volume 내 emptyDir 검색

# emptyDir configuration example 👄

apiVersion: v1 kind: Pod metadata:

name: test-pd

spec:

containers:

- image: registry.k8s.io/test-webserver

name: test-container

volumeMounts:

- mountPath: /cache name: cache-volume

volumes:

- name: cache-volume

emptyDir:

sizeLimit: 500Mi

weblog.yaml 확인 및 생성

```
guru@k8s-master:~$ kubectl run weblog --image=nginx:1.17 --dry-run=client -o yaml
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: null
  labels:
    run: weblog
  name: weblog
spec:
  containers:
  - image: nginx:1.17
   name: weblog
    resources: {}
  dnsPolicy: ClusterFirst
  restartPolicy: Always
status: {}
guru@k8s-master:~$ kubectl run weblog --image=nginx:1.17 --dry-run=client -o yaml > weblog.yaml
```

## vi weblog.yaml 수정

```
● 1 k8s-master ×
apiVersion: v1
kind: Pod
metadata:
 name: weblog
spec:
  containers:
  - image: nginx:1.17
    name: web
    volumeMounts:
    - mountPath: /var/log/nginx
      name: weblog
  - image: busybox
    name: log
    arg: [/bin/sh, -c, "tail -n+1 -f /data/access.lo
    volumeMounts:
    - mountPath: /data
      name: weblog
      readOnly: true
  volumes:
  name: weblog
    emptyDir:
```

## yaml 적용

```
guru@k8s-master:~$ kubectl apply -f weblog.yaml
pod/weblog created
guru@k8s-master:~$ [
```

guru@k8s-master:~\$ kubectl get po weblog NAME READY STATUS RESTARTS AGE weblog 0/2 Pending 0 2m45s

## 18. HostPath Volume

- 1. /data/cka/fluentd.yaml 파일을 만들어 새로은 Pod 생성하세요.
  - 신규생성 Pod Name: fluentd, image: fluentd, namespace: default
- 2. 위 조건을 참고하여 다음 조건에 맞게 볼륨마운트를 설정하시오.
- 3. Worker node의 도커 컨테이너 디렉토리 : /var/lib/docker/containers 동일 디렉토리로 pod에 마운트 하시오.
- 4. Worker node의 /var/log 디렉토리를 fluentd Pod에 동일이름의 디렉토리 마운트 하시오.

kubenetis/docs 내 volume(hostpath) 검색



**Documentation Kubernetes Blog Training Partners Community** 

Q volume(hostpath)

About 643 results (0.18 seconds)

## Volumes | Kubernetes

Kubernetes > docs > concepts > storage > volumes

Dec 18, 2024 ... A hostPath volume mounts a file or directory from the host node's filesystem into your Pod. This is not something that most Pods will need, but ...

# hostPath configuration example

Linux node

Windows node

```
# This manifest mounts /data/foo on the host as /foo inside the
# single container that runs within the hostpath-example-linux Pod.
# The mount into the container is read-only.
apiVersion: v1
kind: Pod
metadata:
 name: hostpath-example-linux
spec:
 os: { name: linux }
 nodeSelector:
    kubernetes.io/os: linux
  containers:
  - name: example-container
    image: registry.k8s.io/test-webserver
   volumeMounts:
    - mountPath: /foo
     name: example-volume
     readOnly: true
 volumes:
  - name: example-volume
    # mount /data/foo, but only if that directory already exists
     path: /data/foo # directory location on host
     type: Directory # this field is optional
```

/data/cka/fluentd.yaml 이동 및 vi 생성

```
guru@k8s-master:~$ cd /data/cka
guru@k8s-master:/data/cka$ vi fluentd.yaml
```

```
piVersion: v1
kind: Pod
metadata:
 name: fluentd
spec:
 containers:

    image: fluentd

   name: fluentd
   ports:
   - containerPort: 80
     protocol: TCP
   volumeMounts:
   - mountPath: /var/lib/docker/container
     name: containersdir
    - mountPath: /var/log
     name: logdir
  volumes:

    name: containersdir

   hostPath:
      path: /var/lib/docker/container
  name: logdir
   hostPath:
      path: /var/log
```

## yaml 설정

```
guru@k8s-master:/data/cka$ kubectl apply -f fluentd.yaml
pod/fluentd created
```

## pod 확인

```
guru@k8s-master:/data/cka$ kubectl get po fluentd
NAME READY STATUS RESTARTS AGE
fluentd 0/1 Pending 0 9s
```

## 19. Persistent Volume

- 1. pv001라는 이름으로 size 1Gi, access mode ReadWriteMany를 사용하여 persistent volume을 생성합니다.
- 2. volume type은 hostPath이고 위치는 /tmp/app-config입니다.



Documentation Kubernetes Blog Training Partners Community



About 4,380 results (0.23 seconds)

## Persistent Volumes | Kubernetes

Kubernetes > docs > concepts > storage > persistent-volumes

Jan 3, 2025 ... A PersistentVolume (PV) is a piece of storage in the cluster that has been provisioned by an administrator or dynamically provisioned using ...

persistent volume 내 kind: 검색

# Persistent Volumes 🖘

Each PV contains a spec and status, which is the specification and status of the volume. The name of a PersistentVolume object must be a valid DNS subdomain name.

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: pv0003
spec:
  capacity:
    storage: 5Gi
 volumeMode: Filesystem
  accessModes:
    - ReadWriteOnce
  persistentVolumeReclaimPolicy: Recycle
  storageClassName: slow
 mountOptions:
    - hard
    - nfsvers=4.1
  nfs:
    path: /tmp
    server: 172.17.0.2
```

## vi pv.yaml 생성 후 코드 붙여넣기 및 수정

```
apiVersion: v1
kind: PersistentVolume
metadata:
   name: pv001
spec:
   capacity:
    storage: 1Gi
   accessModes:
    ReadWriteMany
   hostPath:
    path: /tmp/app-config
```

## yaml 적용

```
guru@k8s-master:~$ kubectl apply -f pv.yaml
persistentvolume/pv001 created
```

## pv 확인

```
guru@k8s-master:~$ kubectl get pv pv001

NAME CAPACITY ACCESS MODES RECLAIM POLICY STATUS CLAIM STORAGECLASS REASON AGE
pv001 1Gi RWX Retain Available 31s
```

## describe pv 확인

```
guru@k8s-master:~$ kubectl describe pv pv001
Name:
                  pv001
Labels:
                  <none>
Annotations:
                  <none>
Finalizers:
                  [kubernetes.io/pv-protection]
StorageClass:
Status:
                  Available
Claim:
Reclaim Policy:
                  Retain
Access Modes:
                  RWX
VolumeMode:
                  Filesystem
Capacity:
                  1Gi
Node Affinity:
                  <none>
Message:
Source:
                    HostPath (bare host directory volume)
    Type:
    Path:
                    /tmp/app-config
    HostPathType:
Events:
                    <none>
guru@k8s-master:~$
```